

```
In [7]: !pip install pandas  
!pip install numpy  
!pip install seaborn  
!pip install matplotlib  
!pip install xgboost  
!pip install scikit-learn
```

Requirement already satisfied: pandas in ./venv/lib/python3.10/site-packages (2.2.2)
Requirement already satisfied: numpy>=1.22.4 in ./venv/lib/python3.10/site-packages (from pandas) (2.1.0)
Requirement already satisfied: python-dateutil>=2.8.2 in ./venv/lib/python3.10/site-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in ./venv/lib/python3.10/site-packages (from pandas) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in ./venv/lib/python3.10/site-packages (from pandas) (2024.1)
Requirement already satisfied: six>=1.5 in ./venv/lib/python3.10/site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
Requirement already satisfied: numpy in ./venv/lib/python3.10/site-packages (2.1.0)
Requirement already satisfied: seaborn in ./venv/lib/python3.10/site-packages (0.13.2)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in ./venv/lib/python3.10/site-packages (from seaborn) (2.1.0)
Requirement already satisfied: pandas>=1.2 in ./venv/lib/python3.10/site-packages (from seaborn) (2.2.2)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in ./venv/lib/python3.10/site-packages (from seaborn) (3.9.2)
Requirement already satisfied: contourpy>=1.0.1 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.0)
Requirement already satisfied: cycler>=0.10 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.53.1)
Requirement already satisfied: kiwisolver>=1.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.5)
Requirement already satisfied: packaging>=20.0 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (24.1)
Requirement already satisfied: pillow>=8 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.1.4)
Requirement already satisfied: python-dateutil>=2.7 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in ./venv/lib/python3.10/site-packages (from pandas>=1.2->seaborn) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in ./venv/lib/python3.10/site-packages (from pandas>=1.2->seaborn) (2024.1)
Requirement already satisfied: six>=1.5 in ./venv/lib/python3.10/site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.16.0)
Requirement already satisfied: matplotlib in ./venv/lib/python3.10/site-packages (3.9.2)
Requirement already satisfied: contourpy>=1.0.1 in ./venv/lib/python3.10/site-packages (from matplotlib) (1.3.0)
Requirement already satisfied: cycler>=0.10 in ./venv/lib/python3.10/site-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in ./venv/lib/python3.10/site-packages (from matplotlib) (4.53.1)
Requirement already satisfied: kiwisolver>=1.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib) (1.4.5)

```
Requirement already satisfied: numpy>=1.23 in ./venv/lib/python3.10/site-packages (from matplotlib) (2.1.0)
Requirement already satisfied: packaging>=20.0 in ./venv/lib/python3.10/site-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=8 in ./venv/lib/python3.10/site-packages (from matplotlib) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib) (3.1.4)
Requirement already satisfied: python-dateutil>=2.7 in ./venv/lib/python3.10/site-packages (from matplotlib)
(2.9.0.post0)
Requirement already satisfied: six>=1.5 in ./venv/lib/python3.10/site-packages (from python-dateutil>=2.7->matplotl
ib) (1.16.0)
Collecting xgboost
  Downloading xgboost-2.1.1-py3-none-manylinux_2_28_x86_64.whl.metadata (2.1 kB)
Requirement already satisfied: numpy in ./venv/lib/python3.10/site-packages (from xgboost) (2.1.0)
Collecting nvidia-nccl-cu12 (from xgboost)
  Downloading nvidia_nccl_cu12-2.23.4-py3-none-manylinux2014_x86_64.whl.metadata (1.8 kB)
Requirement already satisfied: scipy in ./venv/lib/python3.10/site-packages (from xgboost) (1.14.1)
Downloading xgboost-2.1.1-py3-none-manylinux_2_28_x86_64.whl (153.9 MB)
  153.9/153.9 MB 3.2 MB/s eta 0:00:00m eta 0:00:01[36m0:00:02
Downloading nvidia_nccl_cu12-2.23.4-py3-none-manylinux2014_x86_64.whl (199.0 MB)
  199.0/199.0 MB 3.4 MB/s eta 0:00:00m eta 0:00:01[36m0:00:02
Installing collected packages: nvidia-nccl-cu12, xgboost
Successfully installed nvidia-nccl-cu12-2.23.4 xgboost-2.1.1
Requirement already satisfied: scikit-learn in ./venv/lib/python3.10/site-packages (1.5.1)
Requirement already satisfied: numpy>=1.19.5 in ./venv/lib/python3.10/site-packages (from scikit-learn) (2.1.0)
Requirement already satisfied: scipy>=1.6.0 in ./venv/lib/python3.10/site-packages (from scikit-learn) (1.14.1)
Requirement already satisfied: joblib>=1.2.0 in ./venv/lib/python3.10/site-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in ./venv/lib/python3.10/site-packages (from scikit-learn)
(3.5.0)
```

Machine Learning Project: Bank Marketing Dataset

Loading Packages

```
In [8]: import os
```

```
In [9]: import numpy as np
import pandas as pd

import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
In [10]: from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV, StratifiedKFold

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
from sklearn.metrics import confusion_matrix

from sklearn.preprocessing import LabelEncoder, StandardScaler
```

```
In [11]: sns.set_theme(context='paper', style="darkgrid")
```

Loading Dataset

```
In [13]: # Load the dataset
dir_path = './bank+marketing/bank-additional'
file_name = 'bank-additional-full.csv'
# file_name = 'bank-additional.csv'
file_path = os.path.join(dir_path, file_name)
print(f"Dataset: {file_path}")
data = pd.read_csv(file_path, sep=';')
data.head()
```

Dataset: ./bank+marketing/bank-additional/bank-additional-full.csv

```
Out[13]:
```

| | age | job | marital | education | default | housing | loan | contact | month | day_of_week | ... | campaign | pdays | previous | po |
|---|-----|-----------|---------|-------------|---------|---------|------|-----------|-------|-------------|-----|----------|-------|----------|------|
| 0 | 56 | housemaid | married | basic.4y | no | no | no | telephone | may | mon | ... | 1 | 999 | 0 | none |
| 1 | 57 | services | married | high.school | unknown | no | no | telephone | may | mon | ... | 1 | 999 | 0 | none |
| 2 | 37 | services | married | high.school | no | yes | no | telephone | may | mon | ... | 1 | 999 | 0 | none |
| 3 | 40 | admin. | married | basic.6y | no | no | no | telephone | may | mon | ... | 1 | 999 | 0 | none |
| 4 | 56 | services | married | high.school | no | no | yes | telephone | may | mon | ... | 1 | 999 | 0 | none |

5 rows × 21 columns

Dataset Analysis

Initial analysis

```
In [14]: print(data.columns)
```

```
Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',  
      'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',  
      'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',  
      'cons.conf.idx', 'euribor3m', 'nr.employed', 'y'],  
      dtype='object')
```

Dataset Summary

```
In [15]: print("Summary of dataset:")  
print(data.info())
```

Summary of dataset:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 41188 entries, 0 to 41187

Data columns (total 21 columns):

| # | Column | Non-Null Count | Dtype |
|----|----------------|----------------|---------|
| 0 | age | 41188 non-null | int64 |
| 1 | job | 41188 non-null | object |
| 2 | marital | 41188 non-null | object |
| 3 | education | 41188 non-null | object |
| 4 | default | 41188 non-null | object |
| 5 | housing | 41188 non-null | object |
| 6 | loan | 41188 non-null | object |
| 7 | contact | 41188 non-null | object |
| 8 | month | 41188 non-null | object |
| 9 | day_of_week | 41188 non-null | object |
| 10 | duration | 41188 non-null | int64 |
| 11 | campaign | 41188 non-null | int64 |
| 12 | pdays | 41188 non-null | int64 |
| 13 | previous | 41188 non-null | int64 |
| 14 | poutcome | 41188 non-null | object |
| 15 | emp.var.rate | 41188 non-null | float64 |
| 16 | cons.price.idx | 41188 non-null | float64 |
| 17 | cons.conf.idx | 41188 non-null | float64 |
| 18 | euribor3m | 41188 non-null | float64 |
| 19 | nr.employed | 41188 non-null | float64 |
| 20 | y | 41188 non-null | object |

dtypes: float64(5), int64(5), object(11)

memory usage: 6.6+ MB

None

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 41188 entries, 0 to 41187

Data columns (total 21 columns):

| # | Column | Non-Null Count | Dtype |
|---|-----------|----------------|--------|
| 0 | age | 41188 non-null | int64 |
| 1 | job | 41188 non-null | object |
| 2 | marital | 41188 non-null | object |
| 3 | education | 41188 non-null | object |
| 4 | default | 41188 non-null | object |
| 5 | housing | 41188 non-null | object |

```
6   loan                41188 non-null object
7   contact             41188 non-null object
8   month               41188 non-null object
9   day_of_week         41188 non-null object
10  duration            41188 non-null int64
11  campaign            41188 non-null int64
12  pdays               41188 non-null int64
13  previous            41188 non-null int64
14  poutcome            41188 non-null object
15  emp.var.rate        41188 non-null float64
16  cons.price.idx      41188 non-null float64
17  cons.conf.idx       41188 non-null float64
18  euribor3m           41188 non-null float64
19  nr.employed         41188 non-null float64
20  y                   41188 non-null object
dtypes: float64(5), int64(5), object(11)
memory usage: 6.6+ MB
None
```

Missing values in table

```
In [16]: print("Checking for missing values:")
         print(data.isnull().sum())
```

Checking for missing values:

| | |
|----------------|---|
| age | 0 |
| job | 0 |
| marital | 0 |
| education | 0 |
| default | 0 |
| housing | 0 |
| loan | 0 |
| contact | 0 |
| month | 0 |
| day_of_week | 0 |
| duration | 0 |
| campaign | 0 |
| pdays | 0 |
| previous | 0 |
| poutcome | 0 |
| emp.var.rate | 0 |
| cons.price.idx | 0 |
| cons.conf.idx | 0 |
| euribor3m | 0 |
| nr.employed | 0 |
| y | 0 |
| dtype: int64 | |
| age | 0 |
| job | 0 |
| marital | 0 |
| education | 0 |
| default | 0 |
| housing | 0 |
| loan | 0 |
| contact | 0 |
| month | 0 |
| day_of_week | 0 |
| duration | 0 |
| campaign | 0 |
| pdays | 0 |
| previous | 0 |
| poutcome | 0 |
| emp.var.rate | 0 |
| cons.price.idx | 0 |
| cons.conf.idx | 0 |


```
euribor3m      0
nr.employed    0
y              0
dtype: int64
```

Dataset statistics

```
In [17]: print("\nDescriptive statistics:")
data.describe(include='all')
```

Descriptive statistics:

```
Out[17]:
```

| | age | job | marital | education | default | housing | loan | contact | month | day_of_week | ... | campaign |
|---------------|--------------|--------|---------|-------------------|---------|---------|-------|----------|-------|-------------|-----|--------------|
| count | 41188.000000 | 41188 | 41188 | 41188 | 41188 | 41188 | 41188 | 41188 | 41188 | 41188 | ... | 41188.000000 |
| unique | NaN | 12 | 4 | 8 | 3 | 3 | 3 | 2 | 10 | 5 | ... | NaN |
| top | NaN | admin. | married | university.degree | no | yes | no | cellular | may | thu | ... | NaN |
| freq | NaN | 10422 | 24928 | 12168 | 32588 | 21576 | 33950 | 26144 | 13769 | 8623 | ... | NaN |
| mean | 40.02406 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | 2.567593 |
| std | 10.42125 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | 2.770014 |
| min | 17.00000 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | 1.000000 |
| 25% | 32.00000 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | 1.000000 |
| 50% | 38.00000 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | 2.000000 |
| 75% | 47.00000 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | 3.000000 |
| max | 98.00000 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | 56.000000 |

11 rows × 21 columns

Data Visualization

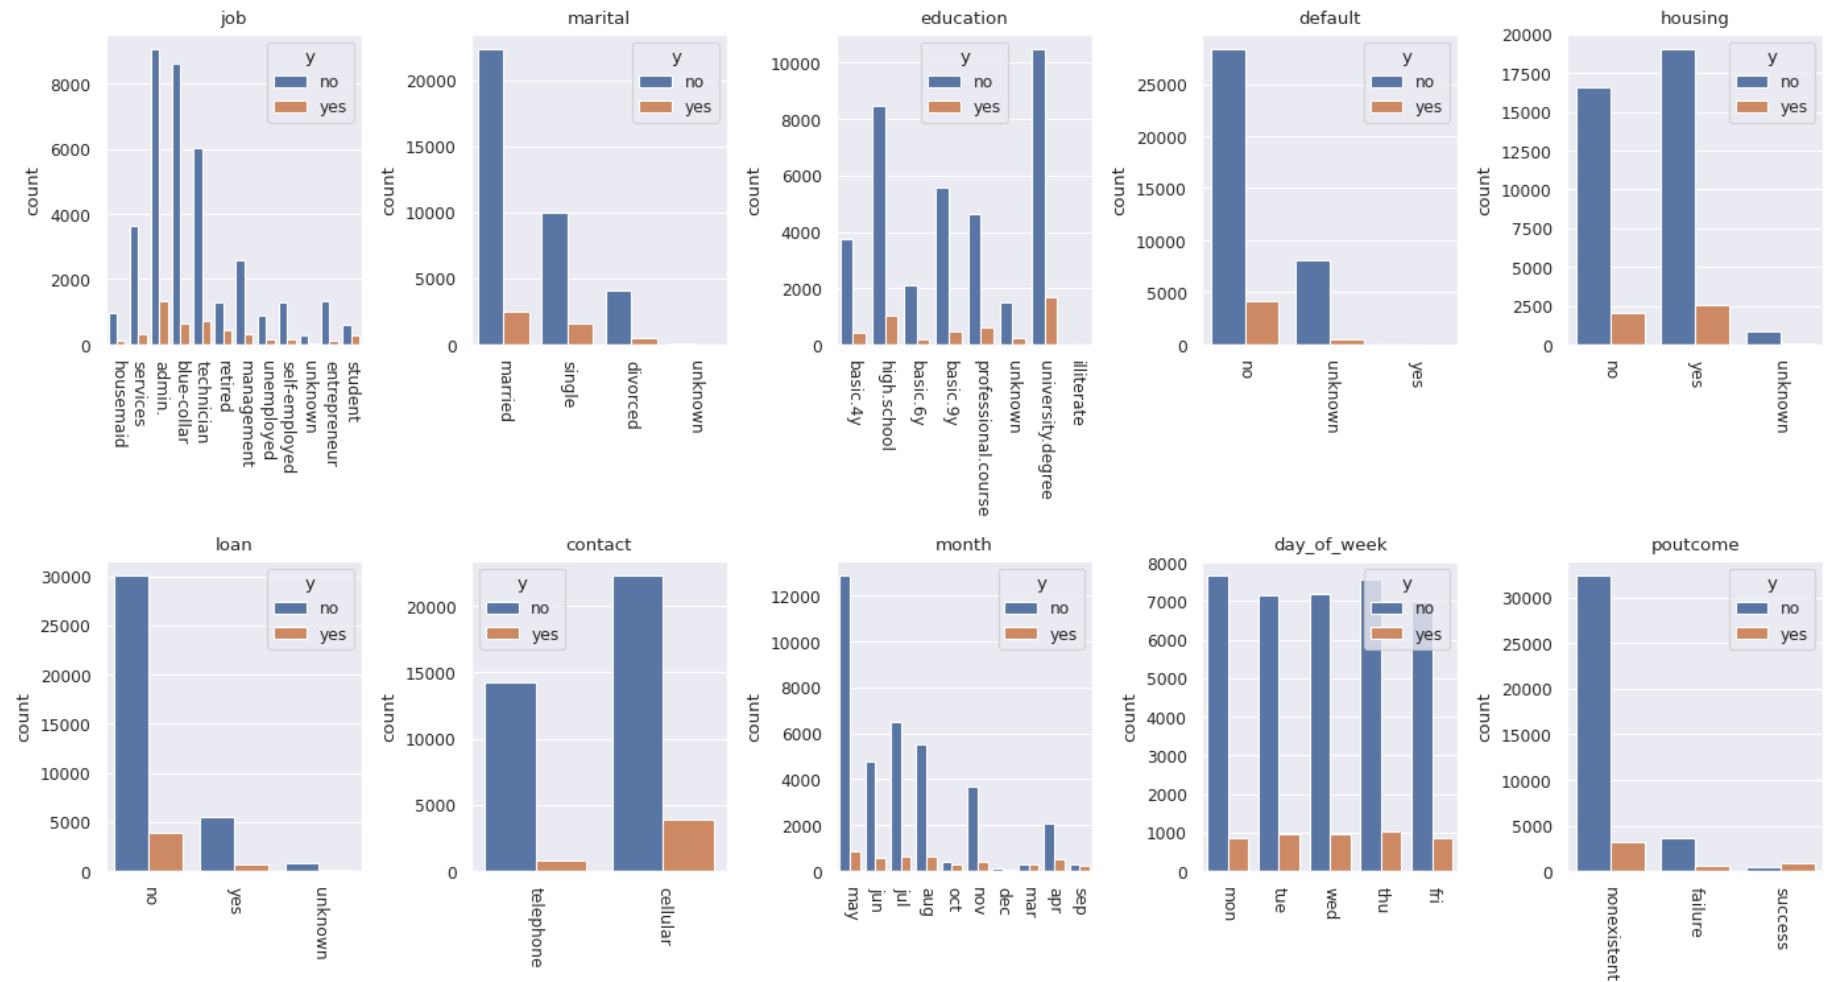
Categorical features against target

```
In [18]: categorical_feature_names = data.drop('y', axis=1).select_dtypes(include=['object']).columns  
print(f"Number of categorical features: {categorical_feature_names.shape[-1]}")
```

Number of categorical features: 10

```
In [19]: fig=plt.figure(figsize=(14, 8))  
for i in range(categorical_feature_names.shape[-1]):  
    plt.subplot(2,5,i+1)  
    sns.countplot(data=data, x=categorical_feature_names[i], hue='y')  
    plt.xticks(rotation=-90)  
    plt.xlabel('')  
    plt.title(categorical_feature_names[i])  
fig.suptitle('Count Plot of Categorical Features', fontsize=16)  
plt.tight_layout()
```

Count Plot of Categorical Features



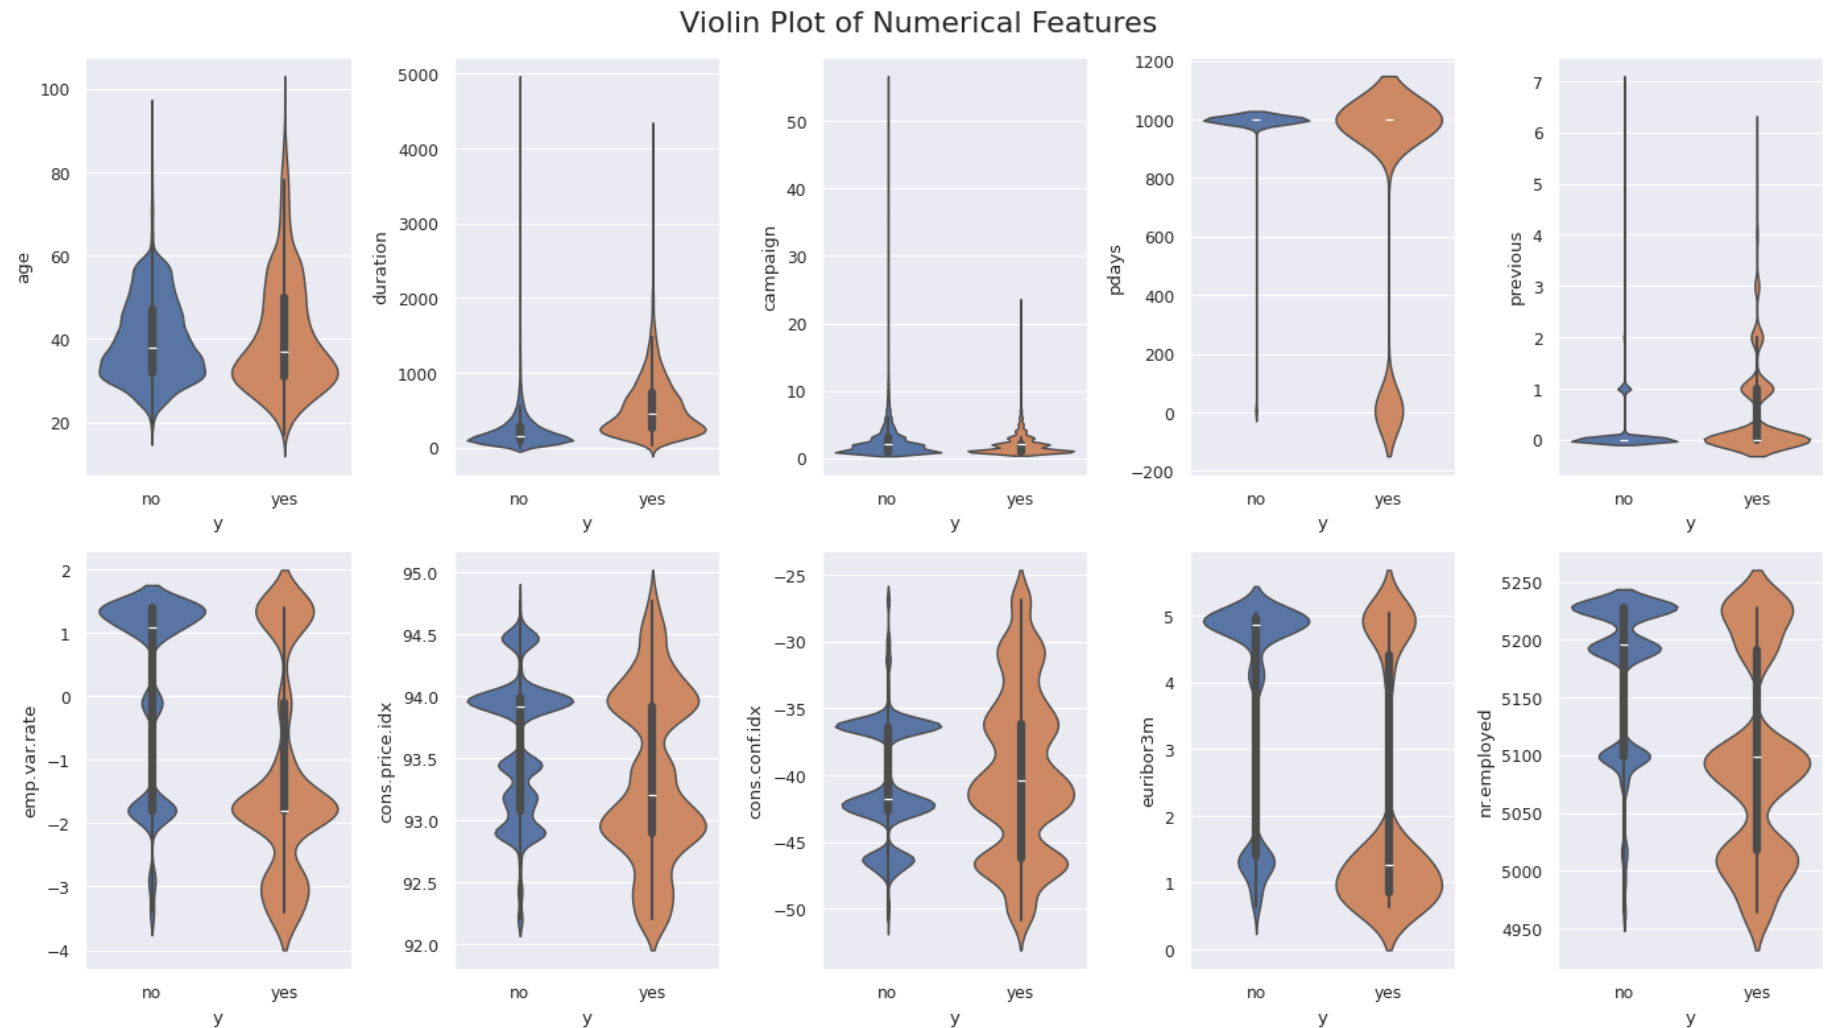
Numerical features against target variable

```
In [20]: numerical_feature_names = data.drop('y', axis=1).select_dtypes(include=['int', 'float']).columns
print(f"Number of numerical features: {numerical_feature_names.shape[-1]}")
```

Number of numerical features: 10

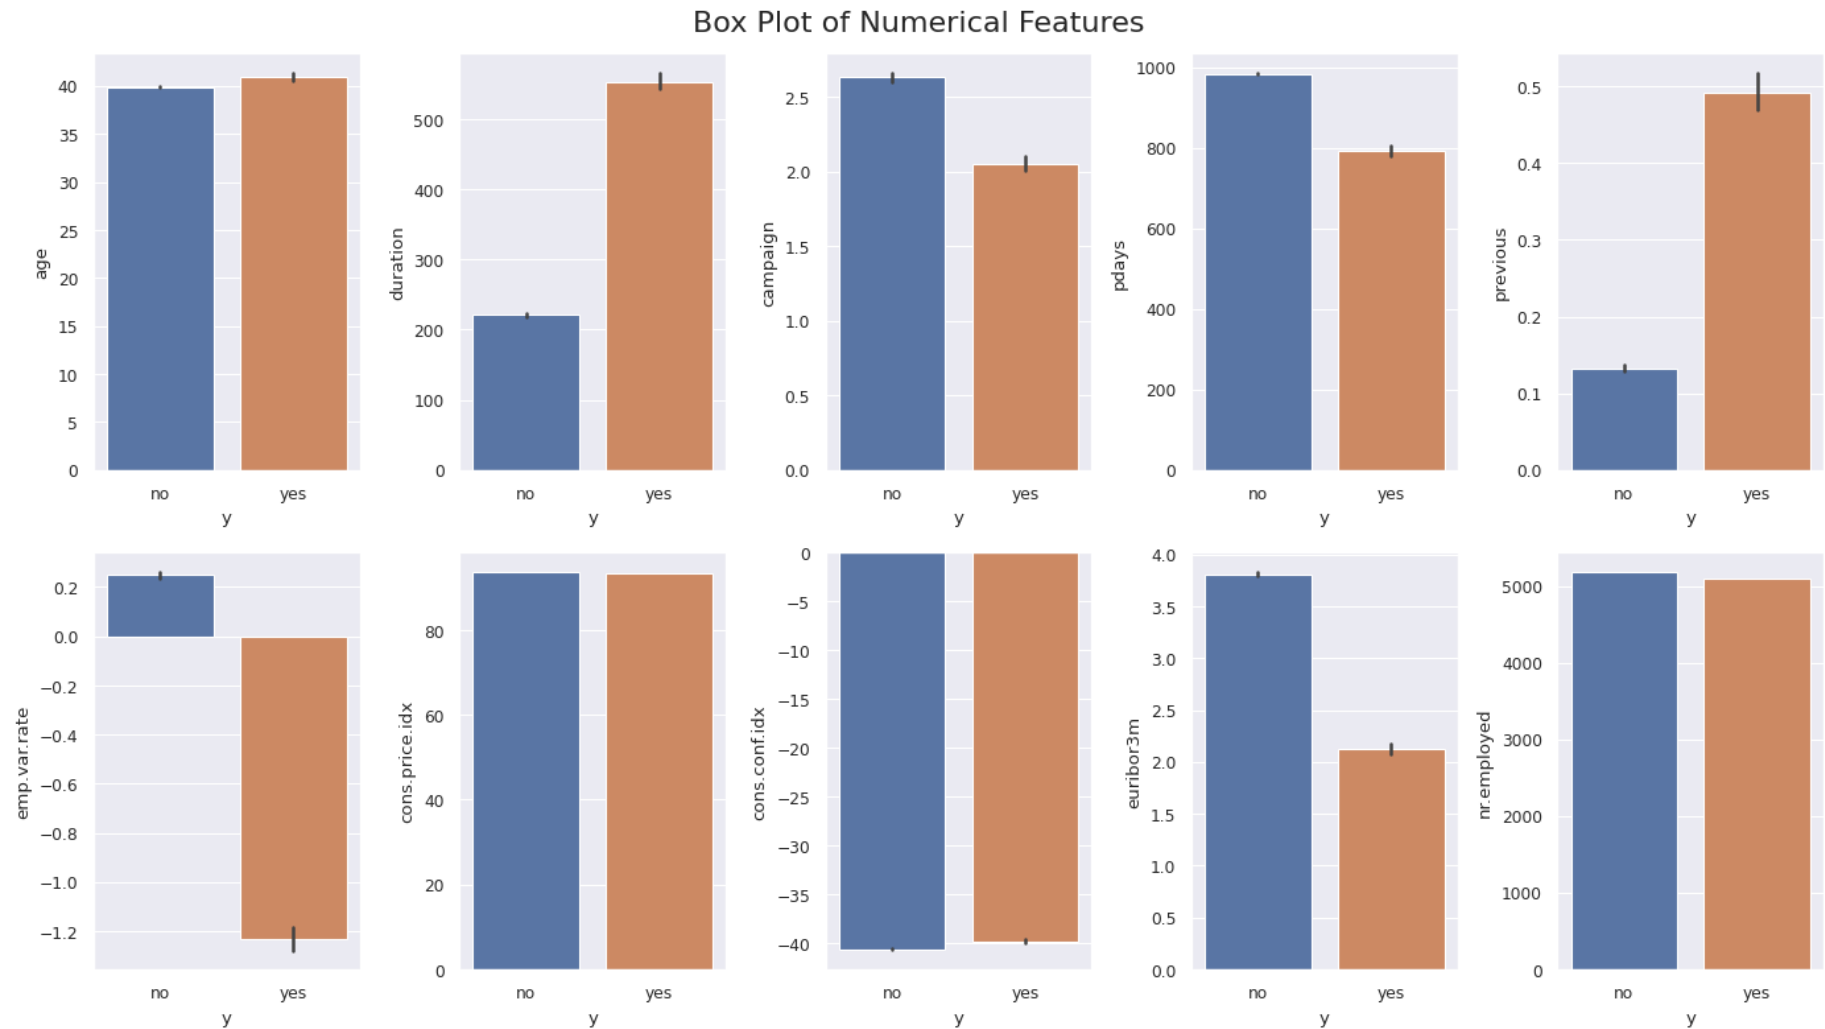
```
In [21]: fig=plt.figure(figsize=(14, 8))
for i in range(numerical_feature_names.shape[-1]):
```

```
plt.subplot(2,5,i+1)
sns.violinplot(data=data, x='y', y=numerical_feature_names[i], hue='y')
fig.suptitle('Violin Plot of Numerical Features', fontsize=16)
plt.tight_layout()
```

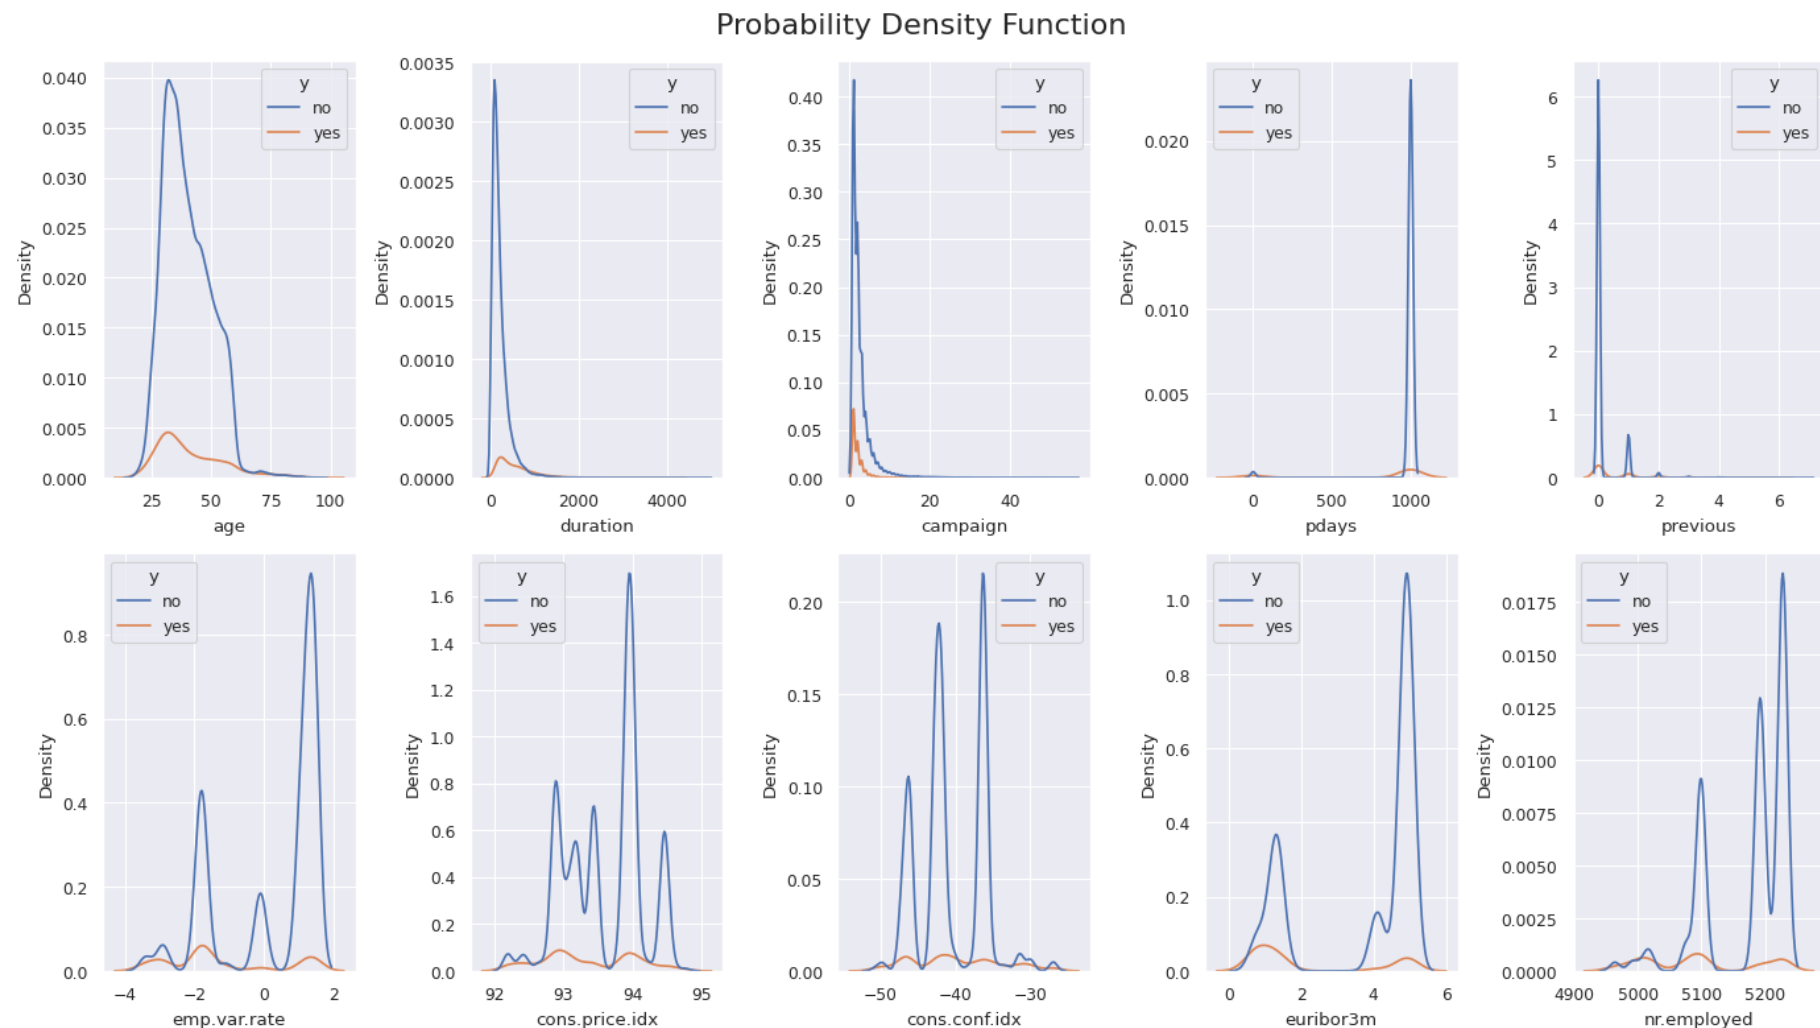


```
In [22]: fig=plt.figure(figsize=(14, 8))
for i in range(numerical_feature_names.shape[-1]):
    plt.subplot(2,5,i+1)
    sns.violinplot(data=data, x='y', y=numerical_feature_names[i], hue='y')
```

```
fig.suptitle('Box Plot of Numerical Features', fontsize=16)
plt.tight_layout()
```



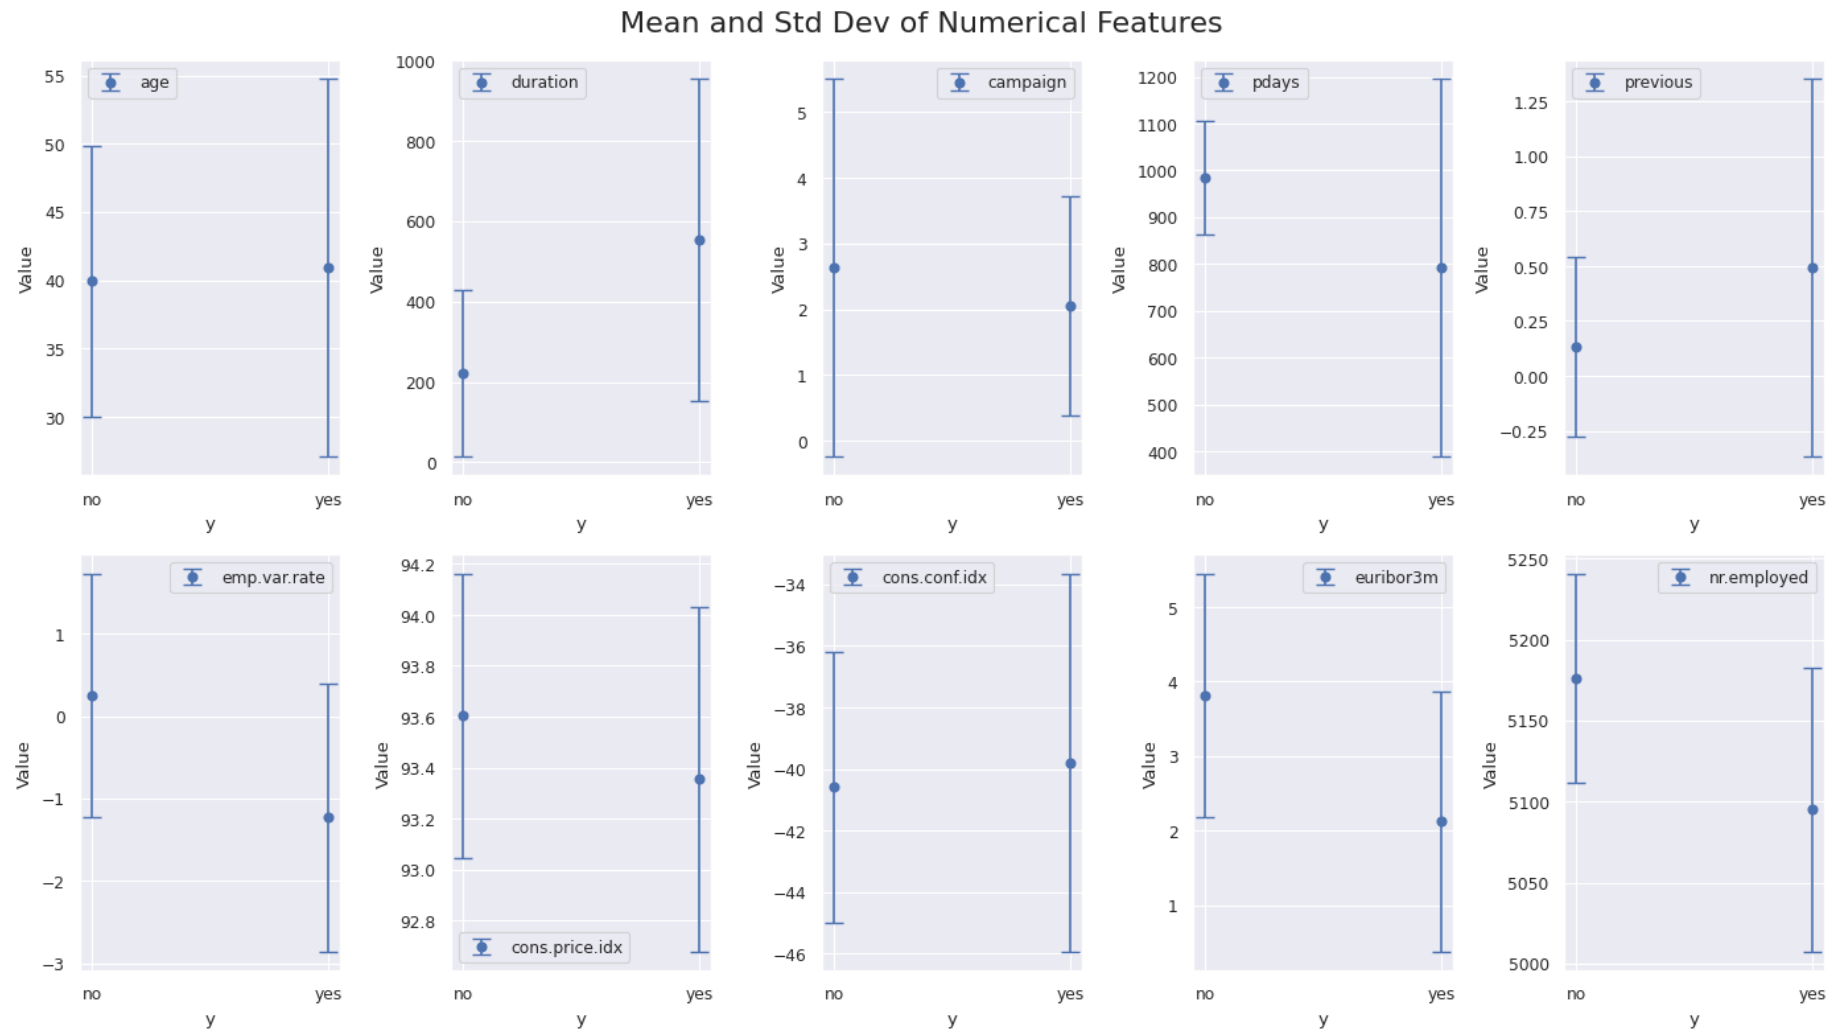
```
In [23]: fig=plt.figure(figsize=(14, 8))
for i in range(numerical_feature_names.shape[-1]):
    plt.subplot(2,5,i+1)
    sns.kdeplot(data=data, x=numerical_feature_names[i], hue='y')
fig.suptitle('Probability Density Function', fontsize=16)
plt.tight_layout()
```



```
In [24]: grouped = data.drop(categorical_feature_names, axis=1).groupby('y').agg(['mean', 'std'])
grouped.columns = ['_'.join(col).strip() for col in grouped.columns.values]
grouped.reset_index(inplace=True)

fig=plt.figure(figsize=(14, 8))
for i in range(numerical_feature_names.shape[-1]):
    plt.subplot(2,5,i+1)
    plt.errorbar(grouped['y'], grouped[f'{numerical_feature_names[i]}_mean'], yerr=grouped[f'{numerical_feature_names[i]}_std'],
    plt.xlabel('y')
```

```
plt.ylabel('Value')
plt.legend()
fig.suptitle('Mean and Std Dev of Numerical Features', fontsize=16)
plt.tight_layout()
```



Class Balance

```
In [25]: plt.figure()
sns.countplot(data=data, x='y', hue='y')
```

```
plt.suptitle("Class Balance")  
plt.tight_layout()
```



Binary Classification

Metrics evaluation function

```
In [26]: def get_metrics(model_name:str, y_pred, y_true):  
         accuracy = accuracy_score(y_true, y_pred)
```



```
precision = precision_score(y_true, y_pred, average='binary')
recall    = recall_score(y_true, y_pred, average='binary')
f1         = f1_score(y_true, y_pred, average='binary')
roc_auc    = roc_auc_score(y_true, y_pred, average="macro")

metrics_table = pd.DataFrame({'model': [model_name],
                              'precision': [precision],
                              'recall': [recall],
                              'f1 score': [f1],
                              'accuracy': [accuracy],
                              'ROC AUC': [roc_auc]})

cm = confusion_matrix(y_pred=y_pred, y_true=y_true)

return metrics_table, cm
```

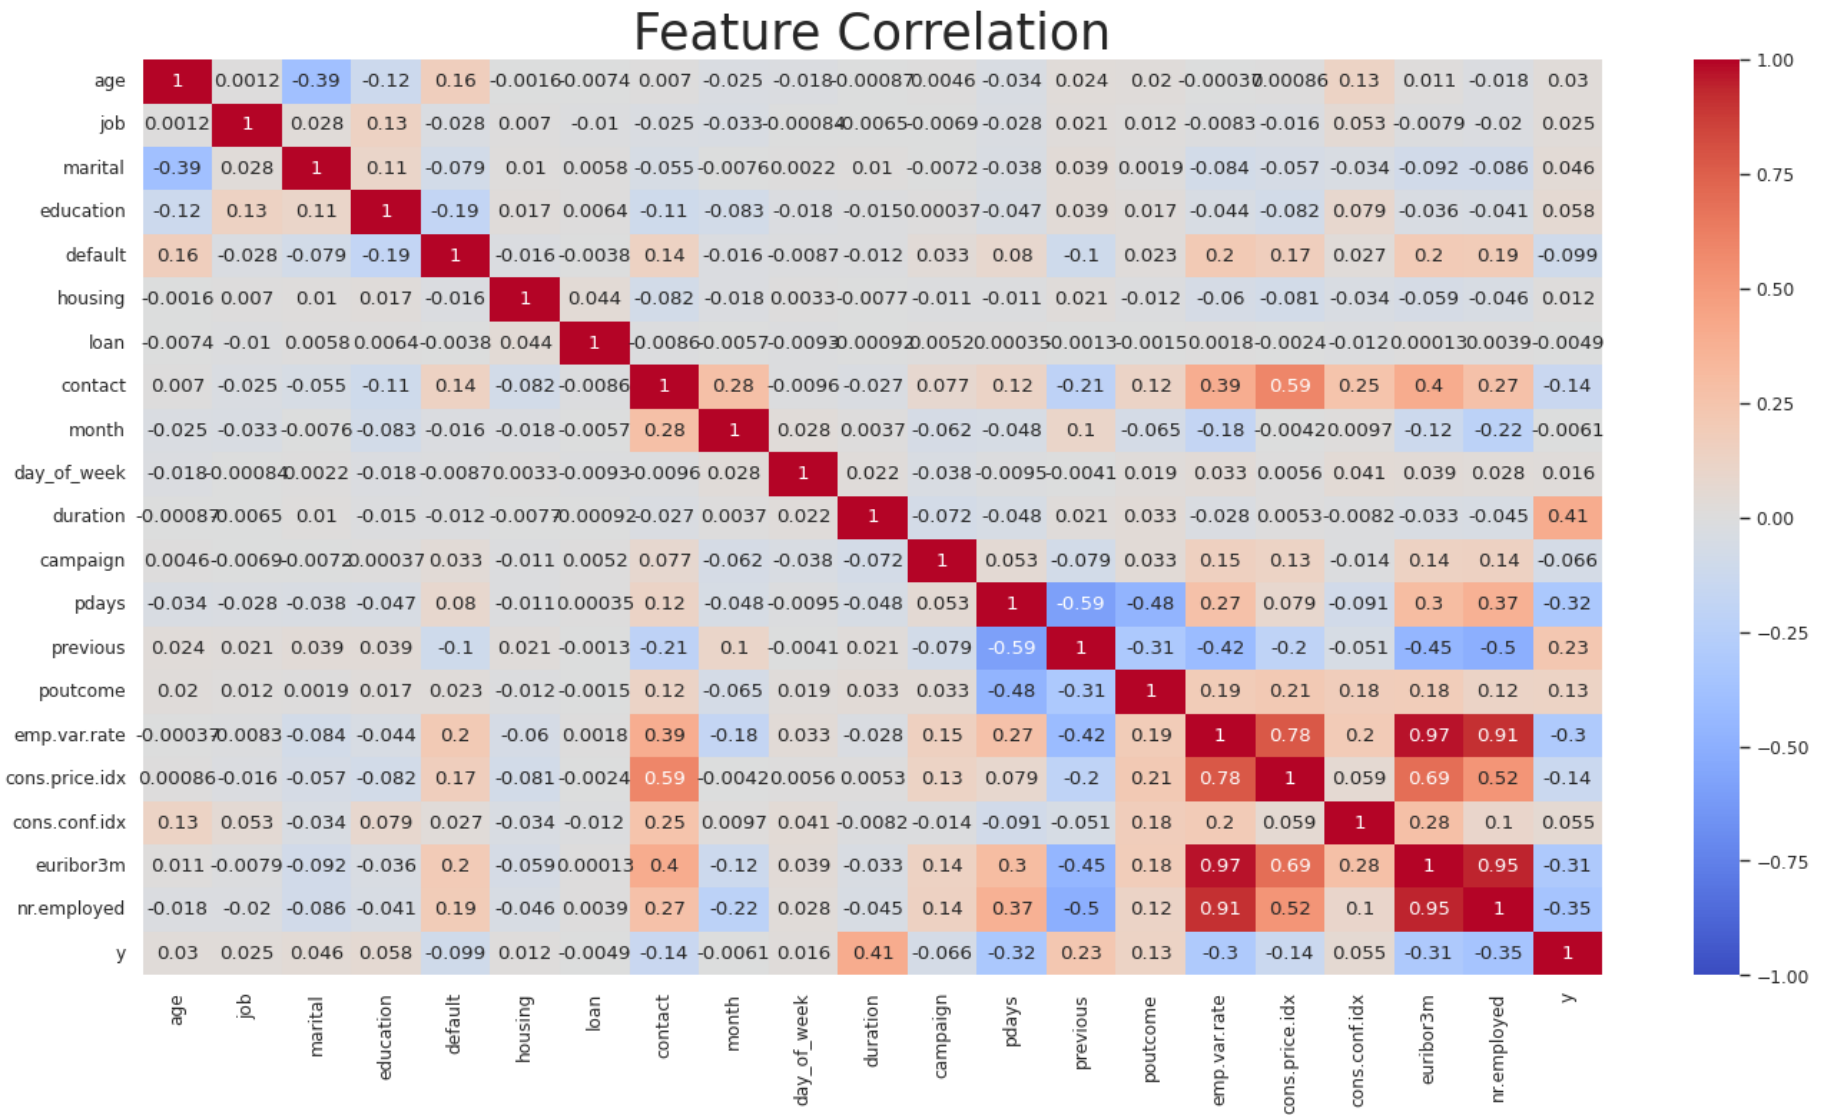
Categorical features encoding

```
In [27]: encoder = LabelEncoder()
for col in categorical_feature_names:
    data[col] = encoder.fit_transform(data[col])

data['y'] = encoder.fit_transform(data['y'])
```

Feature Correlation

```
In [28]: plt.figure(figsize=(14,8))
sns.heatmap(data.corr(method='pearson'), vmin=-1, vmax=1, annot=True, cmap='coolwarm')
plt.title('Feature Correlation', fontsize=25)
plt.tight_layout()
plt.show()
```



```
In [29]: ss_scaler = StandardScaler()
data[categorical_feature_names] = ss_scaler.fit_transform(data[categorical_feature_names])
data.head()
```

```
Out[29]:
```

| | age | job | marital | education | default | housing | loan | contact | month | day_of_week | ... | campaign | pdays | previou |
|---|-----|-----------|-----------|-----------|-----------|-----------|-----------|---------|----------|-------------|-----|----------|-------|---------|
| 0 | 56 | -0.201579 | -0.283741 | -1.753925 | -0.513600 | -1.087707 | -0.452491 | 1.31827 | 0.762558 | -0.718834 | ... | 1 | 999 | |
| 1 | 57 | 0.911227 | -0.283741 | -0.349730 | 1.945327 | -1.087707 | -0.452491 | 1.31827 | 0.762558 | -0.718834 | ... | 1 | 999 | |
| 2 | 37 | 0.911227 | -0.283741 | -0.349730 | -0.513600 | 0.942127 | -0.452491 | 1.31827 | 0.762558 | -0.718834 | ... | 1 | 999 | |
| 3 | 40 | -1.036184 | -0.283741 | -1.285860 | -0.513600 | -1.087707 | -0.452491 | 1.31827 | 0.762558 | -0.718834 | ... | 1 | 999 | |
| 4 | 56 | 0.911227 | -0.283741 | -0.349730 | -0.513600 | -1.087707 | 2.311440 | 1.31827 | 0.762558 | -0.718834 | ... | 1 | 999 | |

5 rows × 21 columns

Train test split

```
In [30]: TEST_SIZE = 0.2
feature_names = data.columns.drop('y')
X = data[feature_names]
y = data['y']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=TEST_SIZE, stratify=y, random_state=42)
```

Simple classifier based on class count

```
In [31]: from collections import Counter

class ClassCountClassifier:
    def fit(self, y):
        class_counts = Counter(y)
        total_samples = len(y)
        self.class_probabilities = {cls: count / total_samples for cls, count in class_counts.items()}
        self.classes = list(self.class_probabilities.keys())
        self.probabilities = list(self.class_probabilities.values())
        return self.class_probabilities

    def predict(self, X):
        # np.random.seed(seed=42)
        return np.random.choice(self.classes, size=len(X), p=self.probabilities)
```

```
In [32]: class_count_clf = ClassCountClassifier()
class_probabilities=class_count_clf.fit(y_train)
print(class_probabilities)
```

```
{0: 0.8873444613050075, 1: 0.11265553869499241}
```

```
In [33]: y_pred = class_count_clf.predict(X_train)
class_count_clf_metrics, class_count_clf_cm = get_metrics("Class Count Classifier", y_pred, y_train)
```

```
In [34]: fig, ax = plt.subplots()
ax=sns.heatmap(class_count_clf_cm,annot=True, ax=ax, fmt='.4g')
ax.xaxis.set_ticklabels(['No Deposit','Deposit'])
ax.yaxis.set_ticklabels(['No Deposit','Deposit'])
ax.set_title('Class Count Classifier-Train',fontsize=18)

print(class_count_clf_metrics.to_string(index=False))
```

| | model | precision | recall | f1 score | accuracy | ROC AUC |
|--|------------------------|-----------|----------|----------|----------|----------|
| | Class Count Classifier | 0.109491 | 0.111261 | 0.110369 | 0.797936 | 0.498188 |



Logistic regression

```
In [37]: lr = LogisticRegression(solver='liblinear', class_weight='balanced', random_state=42)

param_grid = {
    'C': [0.01, 0.1, 1, 10],
    'penalty': ['l1', 'l2']
}
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
grid_search = GridSearchCV(lr, param_grid, cv=cv, scoring='f1', refit=True, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)
lr_clf = grid_search.best_estimator_
print("Best parameters:", grid_search.best_params_)
```


```
print("Best score:", grid_search.best_score_)
```

Fitting 5 folds for each of 8 candidates, totalling 40 fits

[illegible]

```
In [38]: y_pred = lr_clf.predict(X_train)
         lr_clf_metrics, lr_clf_cm = get_metrics("Logistic Regression", y_pred, y_train)
```

| | model | precision | recall | f1 score | accuracy | ROC AUC |
|--|---------------------|-----------|----------|----------|----------|----------|
| | Logistic Regression | 0.428363 | 0.868265 | 0.573692 | 0.854628 | 0.860581 |



Support Vector Machine

```
In [40]: svm = SVC(kernel='rbf', class_weight='balanced', random_state=42)

param_grid = {
    'C': [1, 10, 100],
    'gamma': ['scale', 0.1, 0.01, 0.001],
}

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
grid_search = GridSearchCV(svm, param_grid, cv=cv, scoring='f1', refit=True, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)
svm_clf = grid_search.best_estimator_
```


```
print("Best parameters:", grid_search.best_params_)  
print("Best score:", grid_search.best_score_)
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits
Best parameters: {'C': 1, 'gamma': 0.001}
Best score: 0.573934420125766

```
In [41]: y_pred = svm_clf.predict(X_train)  
         svm_clf_metrics, svm_clf_cm = get_metrics("SVM Classifier", y_pred, y_train)
```

```
In [42]: fig, ax = plt.subplots()  
         ax=sns.heatmap(svm_clf_cm,annot=True, ax=ax, fmt='.4g')  
         ax.xaxis.set_ticklabels(['No Deposit','Deposit'])  
         ax.yaxis.set_ticklabels(['No Deposit','Deposit'])  
         ax.set_title('SVM-Train',fontsize=18)  
  
         print(svm_clf_metrics.to_string(index=False))
```

| | model | precision | recall | f1 score | accuracy | ROC AUC |
|--|----------------|-----------|----------|----------|----------|----------|
| | SVM Classifier | 0.433077 | 0.924838 | 0.589913 | 0.855144 | 0.885567 |



Decision Tree

```
In [43]: dtree = DecisionTreeClassifier(class_weight='balanced', criterion='gini', random_state=42)
param_grid = {
    'max_depth': [5, 10, 20],
    'min_samples_split': [5, 10, 20],
    'min_samples_leaf': [2, 5, 10]
}

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
grid_search = GridSearchCV(dtree, param_grid, cv=cv, scoring='f1', refit=True, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)
dt_clf = grid_search.best_estimator_
```

```
print("Best parameters:", grid_search.best_params_)  
print("Best score:", grid_search.best_score_)
```

Fitting 5 folds for each of 27 candidates, totalling 135 fits

Best parameters: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 5}

Best score: 0.5801383257294856

```
In [44]: y_pred = dt_clf.predict(X_train)  
dt_clf_metrics, dt_clf_cm = get_metrics("Decision Tree Classifier", y_pred, y_train)
```

```
In [45]: fig, ax = plt.subplots()  
ax=sns.heatmap(dt_clf_cm,annot=True, ax=ax, fmt='.4g')  
ax.xaxis.set_ticklabels(['No Deposit', 'Deposit'])  
ax.yaxis.set_ticklabels(['No Deposit', 'Deposit'])  
ax.set_title('Decision Tree-Train', fontsize=18)  
  
print(dt_clf_metrics.to_string(index=False))
```

| | model | precision | recall | f1 score | accuracy | ROC AUC |
|--|--------------------------|-----------|----------|----------|----------|----------|
| | Decision Tree Classifier | 0.472839 | 0.975485 | 0.636939 | 0.874719 | 0.918706 |



Random Forest

```
In [46]: rf = RandomForestClassifier(criterion='gini',  
                                     class_weight='balanced',  
                                     max_features='sqrt',  
                                     oob_score=True,  
                                     bootstrap=True,  
                                     random_state=42)  
  
param_grid = {  
    'n_estimators': [50, 100, 200],  
    'max_depth': [5, 8, 10]  
}
```

```

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
grid_search = GridSearchCV(rf, param_grid, cv=cv, scoring='f1', refit=True, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)
rf_clf = grid_search.best_estimator_
print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)

```

Fitting 5 folds for each of 9 candidates, totalling 45 fits
 Best parameters: {'max_depth': 10, 'n_estimators': 50}
 Best score: 0.6131917036193764

```

In [47]: y_pred = rf_clf.predict(X_train)
         rf_clf_metrics, rf_clf_cm = get_metrics("Random Forest Classifier", y_pred, y_train)

```

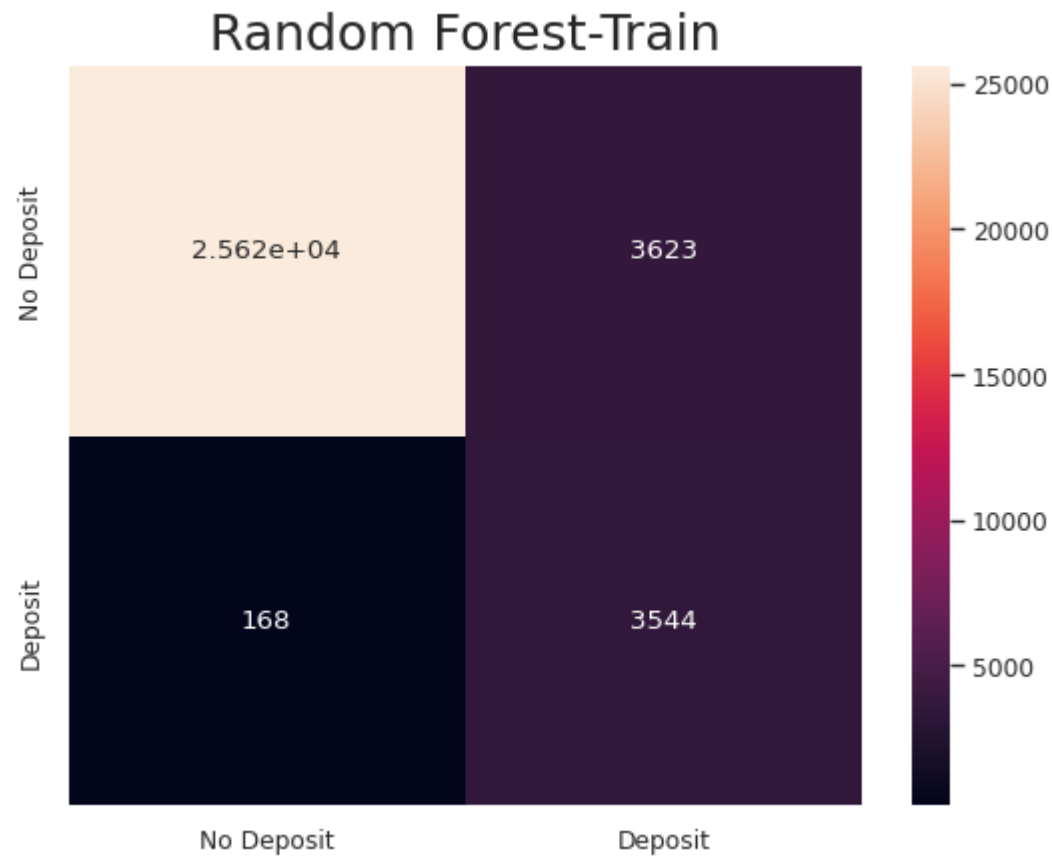
```

In [48]: fig, ax = plt.subplots()
         ax=sns.heatmap(rf_clf_cm,annot=True, ax=ax, fmt='.4g')
         ax.xaxis.set_ticklabels(['No Deposit','Deposit'])
         ax.yaxis.set_ticklabels(['No Deposit','Deposit'])
         ax.set_title('Random Forest-Train',fontsize=18)

         print(rf_clf_metrics.to_string(index=False))
         print(f"Out Of Bag Score (OOB): {rf_clf.oob_score_}")

```

| | model | precision | recall | f1 score | accuracy | ROC AUC |
|--|--------------------------|-----------|----------|----------|----------|----------|
| | Random Forest Classifier | 0.494489 | 0.954741 | 0.65153 | 0.884947 | 0.915414 |
| Out Of Bag Score (OOB): 0.8711684370257967 | | | | | | |



XGBoost

```
In [63]: xgb = XGBClassifier(scale_pos_weight=(y_train == 0).sum() / (y_train == 1).sum(), random_state=42)

param_grid = {'n_estimators': [50, 100, 200],
              'learning_rate': [0.05, 0.1, 0.3], # eta in xgboost documentation - shrinkage
              'max_depth': [4, 6, 8, 10]
             }

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
grid_search = GridSearchCV(xgb, param_grid, cv=cv, scoring='f1', refit=True, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)
xgb_clf = grid_search.best_estimator_
```

```
print("Best parameters:", grid_search.best_params_)  
print("Best score:", grid_search.best_score_)
```

Fitting 5 folds for each of 36 candidates, totalling 180 fits


Best parameters: {'learning_rate': 0.05, 'max_depth': 10, 'n_estimators': 200}

Best score: 0.6302911493186312

```
In [64]: y_pred = xgb_clf.predict(X_train)  
xgb_clf_metrics, xgb_clf_cm = get_metrics("XGBoost", y_pred, y_train)
```

```
In [65]: fig, ax = plt.subplots()  
ax=sns.heatmap(xgb_clf_cm,annot=True, ax=ax, fmt='.4g')  
ax.xaxis.set_ticklabels(['No Deposit', 'Deposit'])  
ax.yaxis.set_ticklabels(['No Deposit', 'Deposit'])  
ax.set_title('XGBoost-Train', fontsize=18)  
  
print(xgb_clf_metrics.to_string(index=False))
```

| model | precision | recall | f1 score | accuracy | ROC AUC |
|---------|-----------|----------|----------|----------|----------|
| XGBoost | 0.652136 | 0.999461 | 0.789278 | 0.939879 | 0.965888 |



Results on Test set

```
In [66]: y_pred = class_count_clf.predict(X_test)
class_count_clf_metrics, class_count_clf_cm = get_metrics("Class Count Classifier", y_pred, y_test)

y_pred = lr_clf.predict(X_test)
lr_clf_metrics, lr_clf_cm = get_metrics("Logistic Regression", y_pred, y_test)

y_pred = svm_clf.predict(X_test)
svm_clf_metrics, svm_clf_cm = get_metrics("SVM Classifier", y_pred, y_test)

y_pred = dt_clf.predict(X_test)
```

```

dt_clf_metrics, dt_clf_cm = get_metrics("Decision Tree Classifier", y_pred, y_test)

y_pred = rf_clf.predict(X_test)
rf_clf_metrics, rf_clf_cm = get_metrics("Random Forest Classifier", y_pred, y_test)

y_pred = xgb_clf.predict(X_test)
xgb_clf_metrics, xgb_clf_cm = get_metrics("XGBoost", y_pred, y_test)

```

```

In [68]: results = pd.concat([class_count_clf_metrics,
                             lr_clf_metrics,
                             svm_clf_metrics,
                             dt_clf_metrics,
                             rf_clf_metrics,
                             xgb_clf_metrics],axis=0)

results.reset_index(drop=True)
# print(results.to_string(index=False))

```

```

Out[68]:

```

| | model | precision | recall | f1 score | accuracy | ROC AUC |
|---|--------------------------|-----------|----------|----------|----------|----------|
| 0 | Class Count Classifier | 0.110521 | 0.112069 | 0.111289 | 0.798373 | 0.498784 |
| 1 | Logistic Regression | 0.433194 | 0.894397 | 0.583685 | 0.856276 | 0.872916 |
| 2 | SVM Classifier | 0.431303 | 0.923491 | 0.587993 | 0.854212 | 0.884454 |
| 3 | Decision Tree Classifier | 0.435560 | 0.892241 | 0.585366 | 0.857611 | 0.872728 |
| 4 | Random Forest Classifier | 0.474444 | 0.920259 | 0.626100 | 0.876184 | 0.895423 |
| 5 | XGBoost | 0.517173 | 0.843750 | 0.641278 | 0.893664 | 0.871875 |

```

In [61]: fig, ax = plt.subplots(nrows=2, ncols=3, figsize=(14,8))
ax = ax.ravel()
ax[0]=sns.heatmap(class_count_clf_cm,annot=True, ax=ax[0], fmt='.4g')
ax[0].xaxis.set_ticklabels(['No Deposit','Deposit'])
ax[0].yaxis.set_ticklabels(['No Deposit','Deposit'])
ax[0].set_title('Class Count Classifier',fontsize=18)

ax[1]=sns.heatmap(lr_clf_cm,annot=True, ax=ax[1], fmt='.4g')
ax[1].xaxis.set_ticklabels(['No Deposit','Deposit'])

```

```
ax[1].yaxis.set_ticklabels(['No Deposit', 'Deposit'])
ax[1].set_title('Logistic Regression', fontsize=18)

ax[2]=sns.heatmap(svm_clf_cm, annot=True, ax=ax[2], fmt='.4g')
ax[2].xaxis.set_ticklabels(['No Deposit', 'Deposit'])
ax[2].yaxis.set_ticklabels(['No Deposit', 'Deposit'])
ax[2].set_title('SVM', fontsize=18)

ax[3]=sns.heatmap(dt_clf_cm, annot=True, ax=ax[3], fmt='.4g')
ax[3].xaxis.set_ticklabels(['No Deposit', 'Deposit'])
ax[3].yaxis.set_ticklabels(['No Deposit', 'Deposit'])
ax[3].set_title('Decision Tree', fontsize=18)

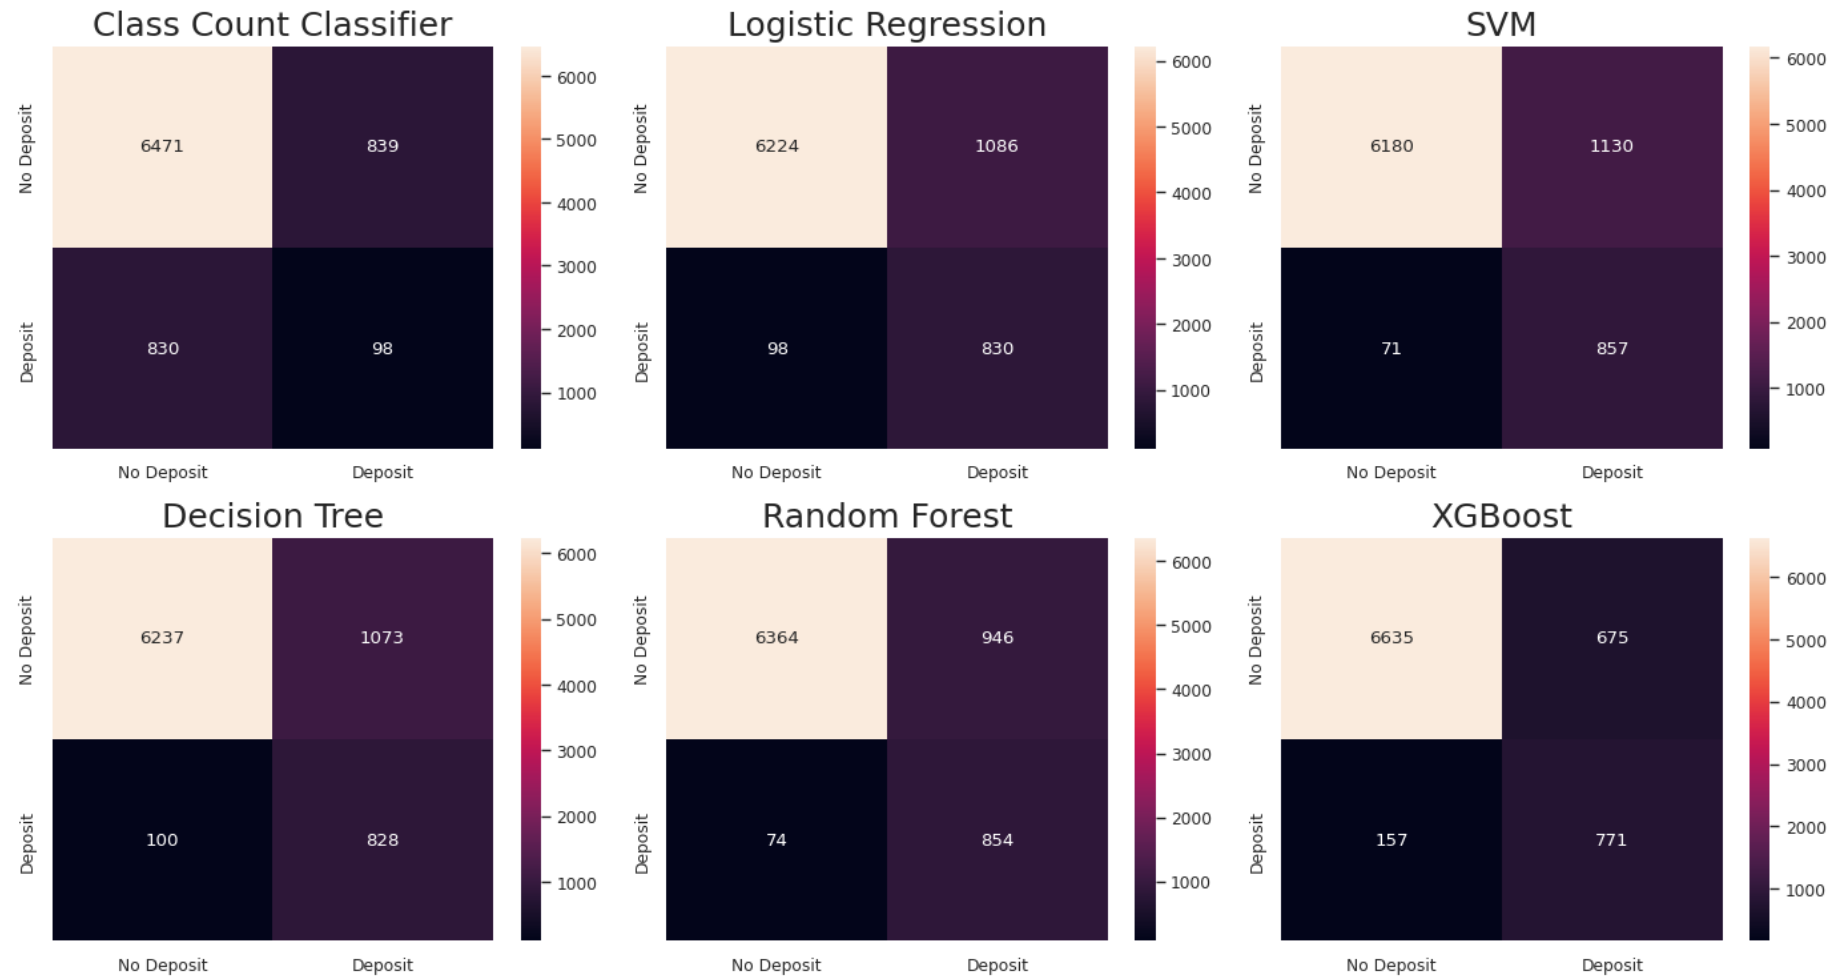
ax[4]=sns.heatmap(rf_clf_cm, annot=True, ax=ax[4], fmt='.4g')
ax[4].xaxis.set_ticklabels(['No Deposit', 'Deposit'])
ax[4].yaxis.set_ticklabels(['No Deposit', 'Deposit'])
ax[4].set_title('Random Forest', fontsize=18)

ax[5]=sns.heatmap(xgb_clf_cm, annot=True, ax=ax[5], fmt='.4g')
ax[5].xaxis.set_ticklabels(['No Deposit', 'Deposit'])
ax[5].yaxis.set_ticklabels(['No Deposit', 'Deposit'])
ax[5].set_title('XGBoost', fontsize=18)

fig.suptitle("Confusion Matrices on Test Set", fontsize=24)

plt.tight_layout()
```

Confusion Matrices on Test Set



Feature Importance

```
In [70]: importance_rf = pd.DataFrame({
          'Features': feature_names,
          'Importance': rf_clf.feature_importances_
        })
importance_xgb = pd.DataFrame({
```

```

    'Features': feature_names,
    'Importance': xgb_clf.feature_importances_
})

importance_rf = importance_rf.sort_values(by='Importance', ascending=False)
importance_xgb = importance_xgb.sort_values(by='Importance', ascending=False)

fig, ax = plt.subplots(ncols=2, figsize=(14,6))
ax[0]=sns.barplot(x='Importance', y='Features', data=importance_rf, ax=ax[0])
ax[0].set_title('Random Forest')

ax[1]=sns.barplot(x='Importance', y='Features', data=importance_xgb, ax=ax[1])
ax[1].set_title('XGBoost')

plt.suptitle('Feature Importance')
plt.tight_layout()
plt.show()

```

